

# Info 206: Computing

Lecture 5

Graphs

September 17, 2015

# Motivating questions


- How does Google compute page rank?
- How does the Internet decide where to move packets?

# Graphs

- Graphs are more general than trees
  - Nodes or vertices
  - Edges
- Two kinds of graphs
  - Directed
  - Undirected


# Undirected Graphs

- In an undirected graph, the edges are lines.
- Undirected graphs show a relationship between two nodes.



# Directed Graphs

- In a directed graph, the edges are arrows
- Directed graphs show the flow from one node to another




# Formal Definition

- A graph  $G = (V, E)$  consists of a finite set of vertices,  $V$ , and a finite set of edges  $E$ .
- Each edge is a pair  $(v, w)$  where  $v, w \in V$

# Directed and undirected

- A **directed** graph, or **digraph**, is a graph in which the edges are ordered pairs
  - $(v, w)$  different from  $(w, v)$
- An **undirected** graph is a graph in which the edges are unordered pairs
  - $(v, w)$  same as  $(w, v)$


# Terminology



- In the undirected graph above, a and b are **adjacent** because  $(a,b) \in E$ . a and b are called **neighbors**.




# Terminology




- In the directed graph above, b is **adjacent** to a because  $(a, b) \in E$ . Note that a is *not* adjacent to b.
- A is a **predecessor** of node B
- B is a **successor** of node A
- The **source** of the edge is node A, the **target** is node B

# Terminology





- A **path** is a sequence of vertices  $w_1, w_2, \dots, w_n$  such that  $(w_i, w_{i+1}) \in E$ ,  $1 \leq i < n$ , and each vertex is unique except that the path may start and end on the same vertex
- The **length** of the path is the number of edges along the path

# Terminology



- An **acyclic path** is a path where each vertex is unique
- A **cyclic path** is a path such that
  - There are at least two vertices on the path
  - $w_1 = w_n$  (path starts and ends at same vertex)


# Cyclic or Acyclic?





# Terminology

- A directed graph that has *no* cyclic paths is called a **DAG** (a Directed Acyclic Graph).
- An undirected graph that has an edge between every pair of vertices is called a **complete** graph.



**Note:** A directed graph can also be a complete graph; in that case, there must be an edge from every vertex to every other vertex.



# Complete or not?



# Complete or not?







# Terminology

- An undirected graph is **connected** if a path exists from every vertex to every other vertex
- A directed graph is **strongly connected** if a path exists from every vertex to every other vertex
- A directed graph is **weakly connected** if a path exists from every vertex to every other vertex, disregarding the direction of the edge




# Connected? (Strongly or weakly?)



# Terminology

- A graph is known as a **weighted graph** if a weight or number is associated with each edge.



# Uses for Graphs

- **WWW**
- **Computer network**


# Single Source Shortest Path

- Problem: find shortest path from a source node to one or more target nodes on a weighted graph
- Dijkstra's Algorithm


# Dijkstra's Algorithm

- Start with source vertex = 0, others =  $\infty$
- Find smallest unvisited vertex
  - Compute new shortest distance to all unvisited nodes
  - Mark vertex as visited


# Dijkstra's Algorithm Example



# Dijkstra's Algorithm Example




# Dijkstra's Algorithm Example







# Dijkstra's Algorithm Example



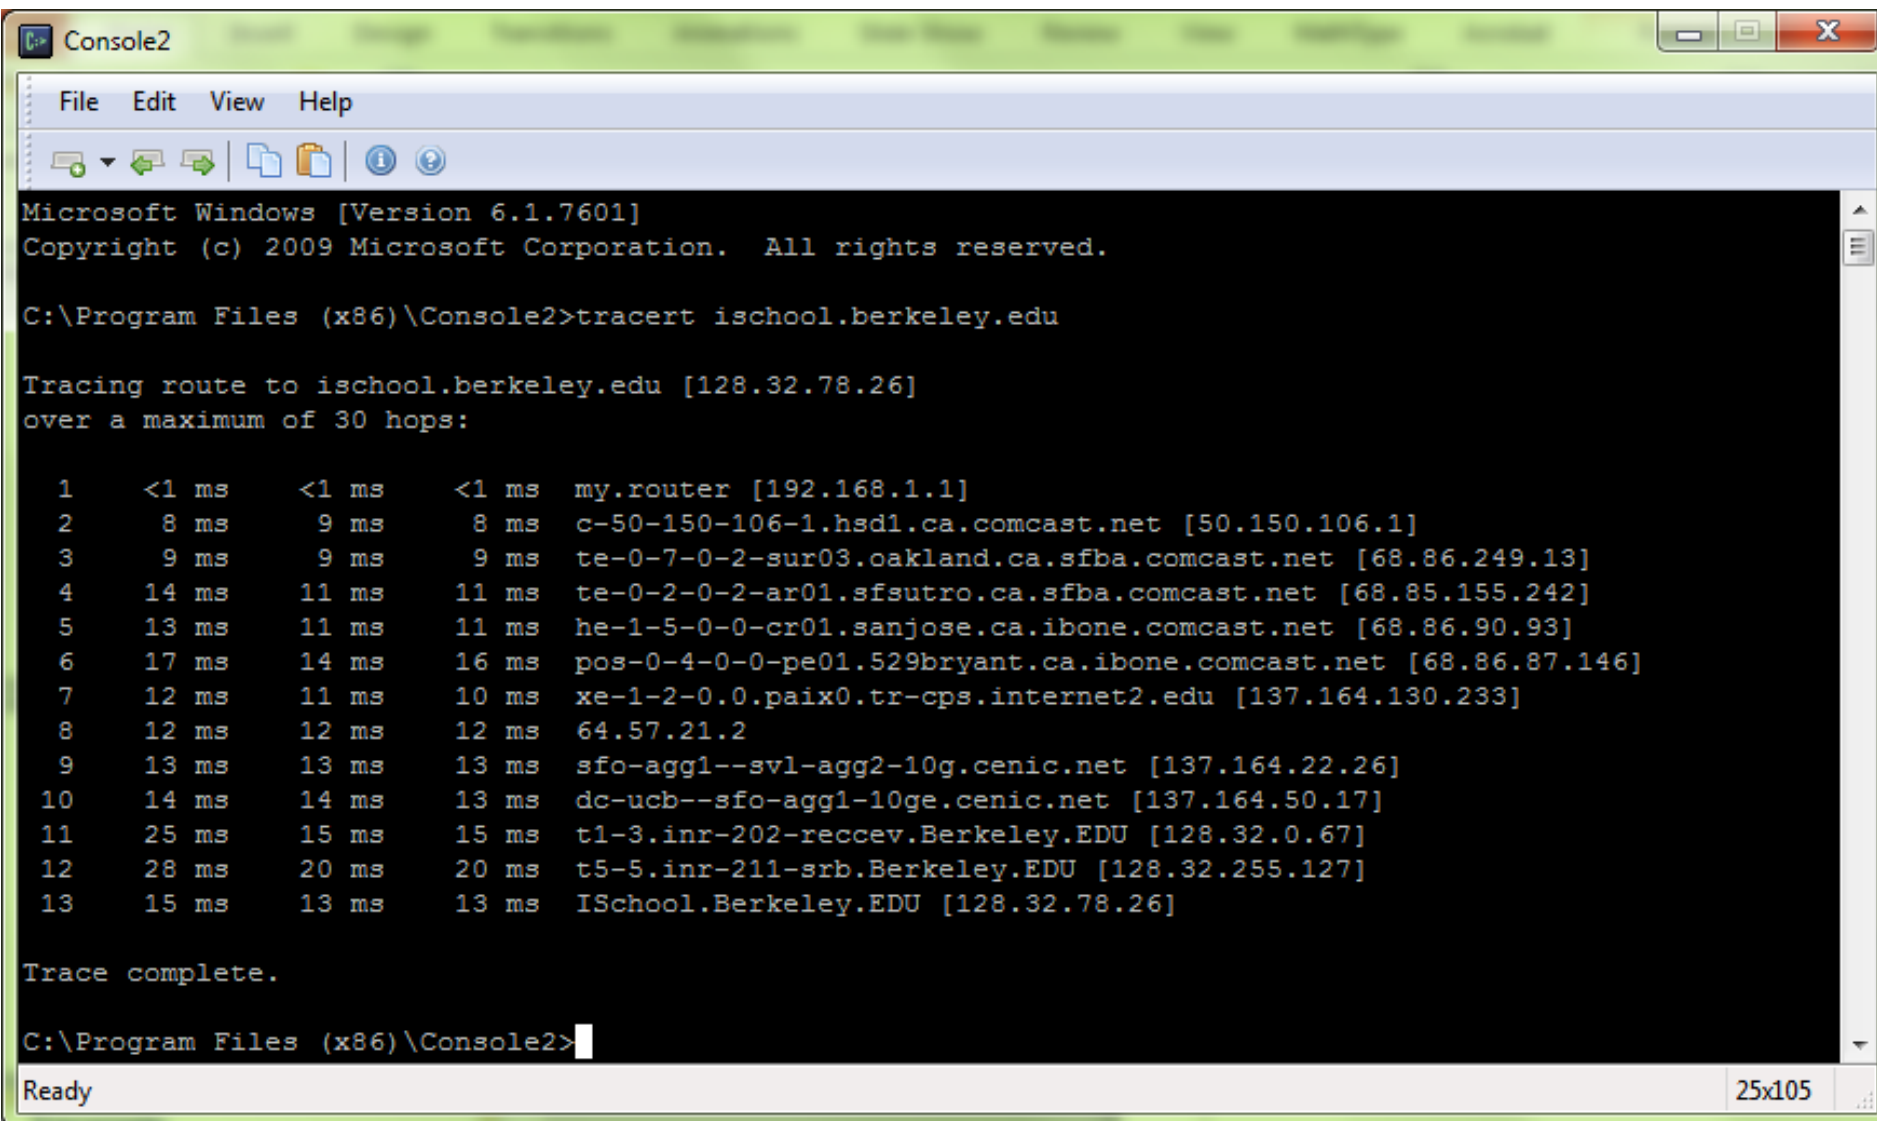
# Dijkstra's Algorithm Example



# Dijkstra's Algorithm Example



# Traceroute



The screenshot shows a Windows console window titled "Console2" with a menu bar (File, Edit, View, Help) and a toolbar. The command prompt shows the execution of the `tracert ischool.berkeley.edu` command. The output displays the path taken by the traceroute, including hop numbers, round-trip times, and IP addresses of the routers. The trace ends at the destination IP `128.32.78.26`.

```
Microsoft Windows [Version 6.1.7601]
Copyright (c) 2009 Microsoft Corporation. All rights reserved.

C:\Program Files (x86)\Console2>tracert ischool.berkeley.edu

Tracing route to ischool.berkeley.edu [128.32.78.26]
over a maximum of 30 hops:

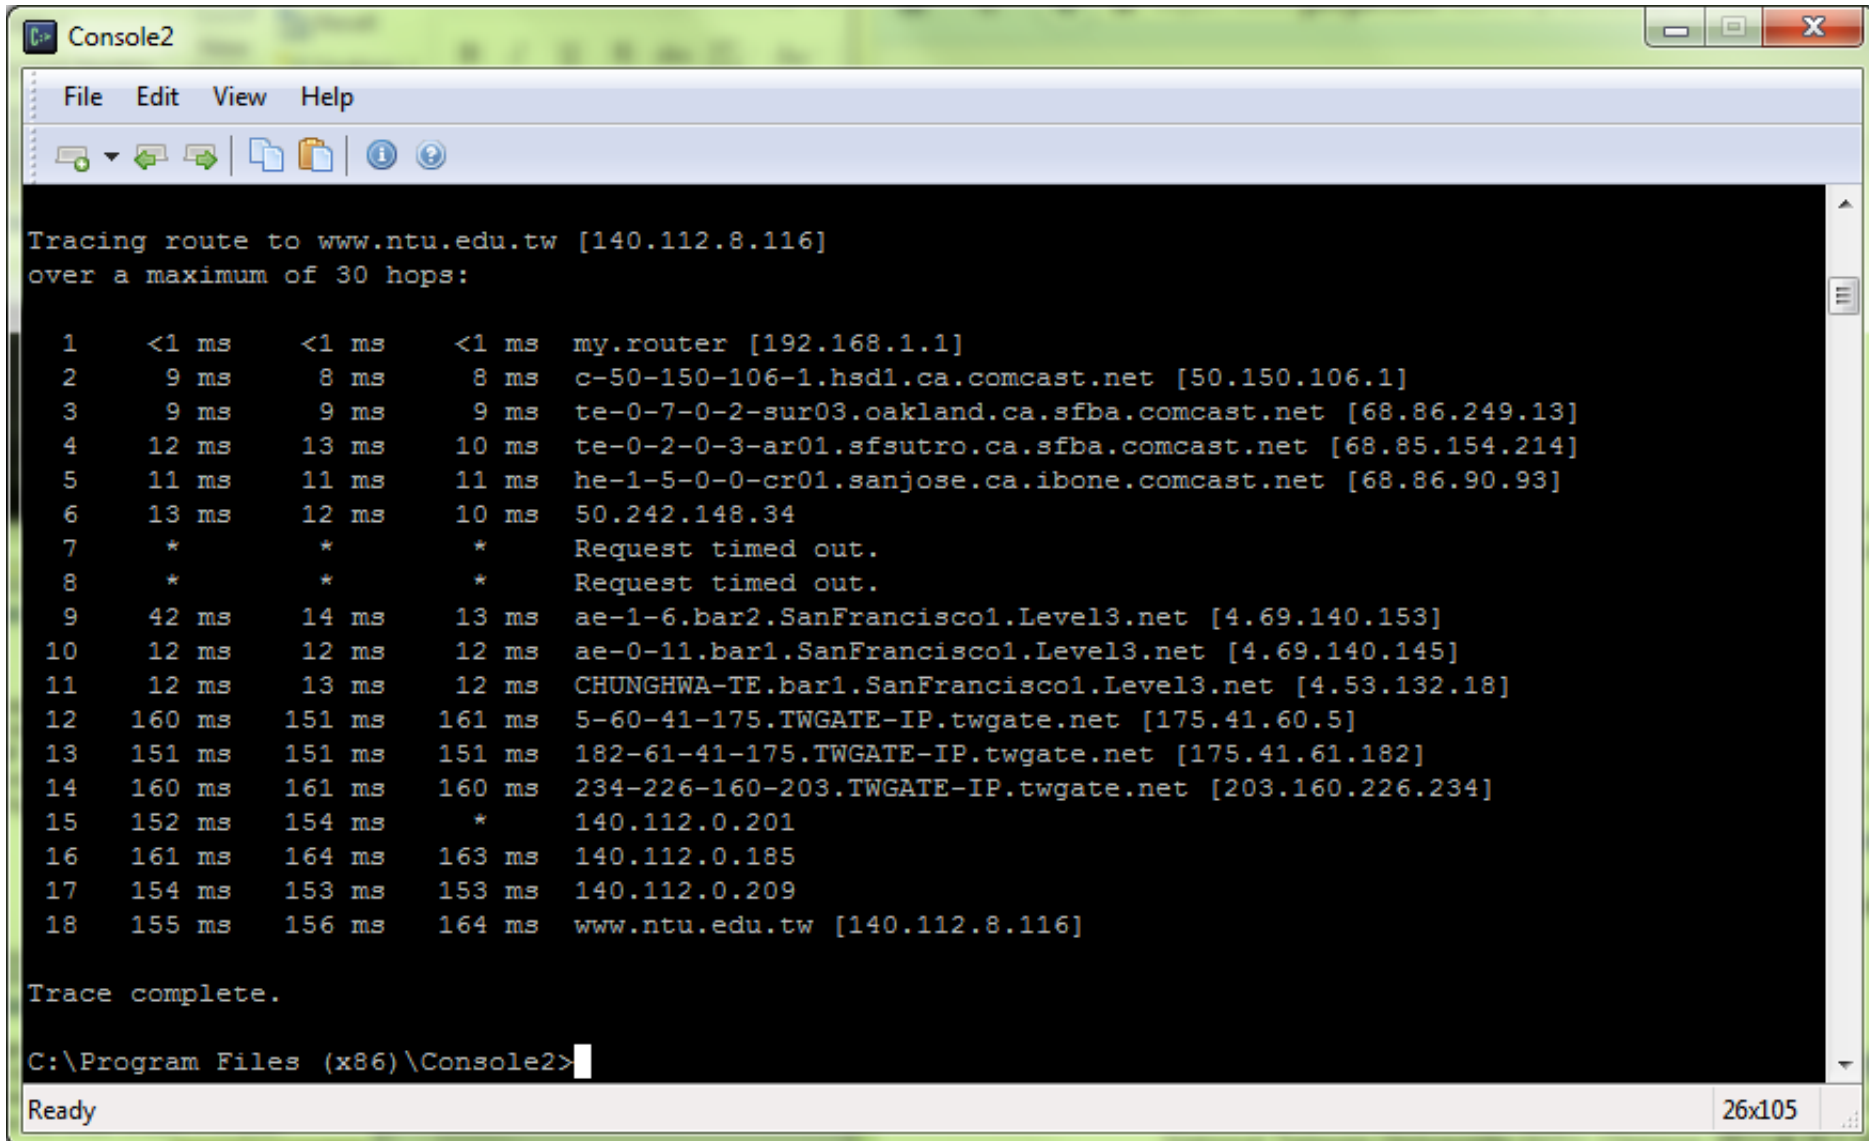
  0  <1 ms    <1 ms    <1 ms    my.router [192.168.1.1]
  1   8 ms     9 ms     8 ms     c-50-150-106-1.hsd1.ca.comcast.net [50.150.106.1]
  2   9 ms     9 ms     9 ms     te-0-7-0-2-sur03.oakland.ca.sfba.comcast.net [68.86.249.13]
  3  14 ms    11 ms    11 ms     te-0-2-0-2-ar01.sfsutro.ca.sfba.comcast.net [68.85.155.242]
  4  13 ms    11 ms    11 ms     he-1-5-0-0-cr01.sanjose.ca.ibone.comcast.net [68.86.90.93]
  5  17 ms    14 ms    16 ms     pos-0-4-0-0-pe01.529bryant.ca.ibone.comcast.net [68.86.87.146]
  6  12 ms    11 ms    10 ms     xe-1-2-0-0.paix0.tr-cps.internet2.edu [137.164.130.233]
  7  12 ms    12 ms    12 ms     64.57.21.2
  8  13 ms    13 ms    13 ms     sfo-agg1--svl-agg2-10g.cenic.net [137.164.22.26]
  9  14 ms    14 ms    13 ms     dc-ucb--sfo-agg1-10ge.cenic.net [137.164.50.17]
 10  25 ms    15 ms    15 ms     t1-3.inr-202-reccev.Berkeley.EDU [128.32.0.67]
 11  28 ms    20 ms    20 ms     t5-5.inr-211-srb.Berkeley.EDU [128.32.255.127]
 12  15 ms    13 ms    13 ms     ISchool.Berkeley.EDU [128.32.78.26]

Trace complete.

C:\Program Files (x86)\Console2>
```

Ready 25x105

# Traceroute



The screenshot shows a Windows console window titled "Console2" with a menu bar (File, Edit, View, Help) and a toolbar. The main text area displays the output of a traceroute command. The command is "Tracing route to www.ntu.edu.tw [140.112.8.116] over a maximum of 30 hops:". The output shows 18 hops. Hops 1-6 show successful connections with round-trip times. Hops 7 and 8 show "Request timed out." Hops 9-18 show successful connections to the destination. The status bar at the bottom shows "Ready" and "26x105".

```
Tracing route to www.ntu.edu.tw [140.112.8.116]
over a maximum of 30 hops:

  1  <1 ms    <1 ms    <1 ms    my.router [192.168.1.1]
  2   9 ms     8 ms     8 ms     c-50-150-106-1.hsd1.ca.comcast.net [50.150.106.1]
  3   9 ms     9 ms     9 ms     te-0-7-0-2-sur03.oakland.ca.sfba.comcast.net [68.86.249.13]
  4  12 ms    13 ms    10 ms    te-0-2-0-3-ar01.sfsutro.ca.sfba.comcast.net [68.85.154.214]
  5  11 ms    11 ms    11 ms    he-1-5-0-0-cr01.sanjose.ca.ibone.comcast.net [68.86.90.93]
  6  13 ms    12 ms    10 ms    50.242.148.34
  7   *        *        *        Request timed out.
  8   *        *        *        Request timed out.
  9  42 ms    14 ms    13 ms    ae-1-6.bar2.SanFrancisco1.Level3.net [4.69.140.153]
 10  12 ms    12 ms    12 ms    ae-0-11.bar1.SanFrancisco1.Level3.net [4.69.140.145]
 11  12 ms    13 ms    12 ms    CHUNGHWA-TE.bar1.SanFrancisco1.Level3.net [4.53.132.18]
 12 160 ms   151 ms   161 ms    5-60-41-175.TWGATE-IP.twgate.net [175.41.60.5]
 13 151 ms   151 ms   151 ms    182-61-41-175.TWGATE-IP.twgate.net [175.41.61.182]
 14 160 ms   161 ms   160 ms    234-226-160-203.TWGATE-IP.twgate.net [203.160.226.234]
 15 152 ms   154 ms    *        140.112.0.201
 16 161 ms   164 ms   163 ms    140.112.0.185
 17 154 ms   153 ms   153 ms    140.112.0.209
 18 155 ms   156 ms   164 ms    www.ntu.edu.tw [140.112.8.116]

Trace complete.

C:\Program Files (x86)\Console2>
```

Ready 26x105

# Random Walks Over the Web


- Model:
  - User starts at a random Web page
  - User randomly clicks on links, surfing from page to page
- How much time is spent on each page?
- This is PageRank (named after Larry Page)

# PageRank: Defined

Given page  $x$  with in-bound links  $t_1 \dots t_n$ , where

- $C(t)$  is the out-degree of  $t$
- $\alpha$  is probability of random jump
- $N$  is the total number of nodes in the graph

$$PR(x) = \alpha \left( \frac{1}{N} \right) + (1 - \alpha) \sum_{i=1}^n \frac{PR(t_i)}{C(t_i)}$$



# PageRank

Page rank of x

Weighted probability of clicking into x from an adjacent page

$$PR(x) = \alpha \left( \frac{1}{N} \right) + (1 - \alpha) \sum_{i=1}^n \frac{PR(t_i)}{C(t_i)}$$

Probability of a random “jumping” to x



# Computing PageRank

- Properties of PageRank
  - Can be computed iteratively
  - Effects at each iteration is local
- Sketch of algorithm:
  - Start with seed  $PR_i$  values
  - Each page distributes  $PR_i$  “credit” to all pages it links to
  - Each target page adds up “credit” from multiple in-bound links to compute  $PR_{i+1}$
  - Iterate until values converge